build_c4_support_eval.py

`scripts/build_c4_support_eval.py · 160 lines`

What it does

Builds the customer-support evaluation set by streaming the C4 validation split (no full download) and keeping documents that pass the support filter, until enough are accepted or the scan limit is reached.

Run

```
uv run python scripts/build_c4_support_eval.py [--preview N]
```

Pipeline order: inspect_model -> build_c4_support_eval -> collect_stats -> score_experts -> replace_experts -> evaluate (before_training) -> train_biases -> evaluate (after_training) -> export_model

Reads

- `allenai/c4` (streamed)
- filter settings `C4_*`

Writes

- `EVALUATION_DATASET_PATH` (JSONL)
- `<same name>.manifest.json` with counts and settings

Notes

- `--preview N` prints the first N accepted documents without writing anything.
- Reports scanned/accepted/rejected/duplicate counts and the most common rejection reasons.

Defined in this file

Extracted from the source code (classes, public functions and methods).

Name	Description
`main()`	Script entry point: runs this stage end to end.

Settings it reads

`C4_CONFIG`, `C4_DATASET_ID`, `C4_DOMAIN`, `C4_MAX_SAMPLES`, `C4_MAX_SCANNED`, `C4_MAX_TEXT_LENGTH`, `C4_MIN_KEYWORD_MATCHES`, `C4_MIN_TEXT_LENGTH`, `C4_SPLIT`, `EVALUATION_DATASET_PATH`, `SEED`

Connections

Uses: `moe_pruning.data.c4_support`, `moe_pruning.utils.cli`, `moe_pruning.utils.paths`