convert_model.py

`scripts/convert_model.py · 65 lines`

What it does

Saves a copy of the base model in `MODEL_DTYPE`, e.g. OLMoE float32 (28 GB) to float16 (14 GB), so it can be stored as a Kaggle Dataset and used via `MODEL_PATH` without repeating the large download.

Run

```
uv run python scripts/convert_model.py --output <dir>
```

Pipeline order: inspect_model -> build_c4_support_eval -> collect_stats -> score_experts -> replace_experts -> evaluate (before_training) -> train_biases -> evaluate (after_training) -> export_model

Reads

- the configured base model

Writes

- `<dir>/` with safetensors shards, tokenizer and `conversion.json`

Notes

- Runs on CPU and needs roughly the target size in RAM.

Defined in this file

Extracted from the source code (classes, public functions and methods).

Name	Description
`main()`	Script entry point: runs this stage end to end.

Settings it reads

`MODEL_DTYPE`, `MODEL_PATH or HF_MODEL_ID`

Connections

Uses: `moe_pruning.model.loader`, `moe_pruning.utils.cli`, `moe_pruning.utils.device`, `moe_pruning.utils.paths`