export_model.py

`scripts/export_model.py · 145 lines`

What it does

Writes the final pruned model to `PRUNED_MODEL_PATH` and verifies it: reloads it with the custom loader, checks that no removed expert weights are stored, and compares outputs with the in-memory model.

Run

```
uv run python scripts/export_model.py [--source trained|replaced] [--format full|delta]
```

Pipeline order: inspect_model -> build_c4_support_eval -> collect_stats -> score_experts -> replace_experts -> evaluate (before_training) -> train_biases -> evaluate (after_training) -> export_model

Reads

- `trained-model/` (bias method) or `replaced-model/`
- experiment reports

Writes

- `PRUNED_MODEL_PATH/` (checkpoint + experiment metadata)
- `export_verification.json`

Defined in this file

Extracted from the source code (classes, public functions and methods).

Name	Description
`main()`	Script entry point: runs this stage end to end.

Settings it reads

`CHECKPOINT_FORMAT`, `PRUNED_MODEL_PATH`, `REPLACEMENT_METHOD`

Connections

Uses: `moe_pruning.data.calibration`, `moe_pruning.model.architecture`, `moe_pruning.model.loader`, `moe_pruning.serialization.checkpoint`, `moe_pruning.utils.cli`, `moe_pruning.utils.device`, `moe_pruning.utils.paths`