score_experts.py

`scripts/score_experts.py · 93 lines`

What it does

Ranks experts per layer by EASY-EP importance and selects the least important ones in the configured layers.

Run

```
uv run python scripts/score_experts.py
```

Pipeline order: inspect_model -> build_c4_support_eval -> collect_stats -> score_experts -> replace_experts -> evaluate (before_training) -> train_biases -> evaluate (after_training) -> export_model

Reads

- `expert_stats.pt` and its manifest

Writes

- `expert_scores.json`, `expert_scores.csv`
- `expert_selection.json`

Notes

- Logs each layer's ranking with each expert's share of the layer's importance.

Defined in this file

Extracted from the source code (classes, public functions and methods).

Name	Description
`main()`	Script entry point: runs this stage end to end.

Settings it reads

`MODEL_PATH or HF_MODEL_ID`, `PRUNE_EXPERTS_PER_LAYER`, `PRUNE_LAYERS`, `PRUNE_UNOBSERVED_POLICY`, `PRUNING_METHOD`

Connections

Uses: `moe_pruning.model.architecture`, `moe_pruning.pruning.expert_selection`, `moe_pruning.scoring.statistics`, `moe_pruning.utils.cli`, `moe_pruning.utils.paths`