metrics.py

`src/moe_pruning/evaluation/metrics.py · 56 lines · module moe_pruning.evaluation.metrics`

What it does

Small, model-independent metric helpers used by the evaluation script.

How it works

- `summarize_losses()`: mean token loss (micro-average), mean example loss, counts, and perplexity when the task makes it meaningful (causal LM only).
- `compare_to_baseline()`: absolute and relative loss change against the original model.
- `per_example_comparison()`: mean, median and max per-example loss change and the share of examples that got worse.
- `prediction_agreement()`: share of positions where the pruned model's top-1 token equals the original model's.

Defined in this file

Extracted from the source code (classes, public functions and methods).

Name	Description
`summarize_losses(loss_sums, token_counts, perplexity)`	Mean token loss, mean example loss, counts (and perplexity for causal LMs).
`compare_to_baseline(stage, baseline)`	Absolute and relative loss change against the original model.
`per_example_comparison(stage_losses, base_losses)`	Mean / median / max per-example loss change and share of examples worse.
`prediction_agreement(stage_preds, base_preds)`	Fraction of target positions whose teacher-forced argmax matches the baseline.

Connections

Used by: `scripts/evaluate.py`, `src/moe_pruning/evaluation/evaluator.py`