expert_selection.py

`src/moe_pruning/pruning/expert_selection.py · 82 lines · module moe_pruning.pruning.expert_selection`

What it does

Turns EASY-EP statistics into a per-layer ranking and selects the experts to replace.

How it works

- Every layer is ranked independently by importance (ties broken by expert id).
- Only layers chosen by `PRUNE_LAYERS` get a selection of the `PRUNE_EXPERTS_PER_LAYER` least important experts.
- `ExpertScore` rows hold importance, its share in the layer, rank, token count, routing frequency, average gate, output norm and token contribution; they become `expert_scores.json/.csv`.

Defined in this file

Extracted from the source code (classes, public functions and methods).

Name	Description
`class ExpertScore`	One row of expert_scores.json/.csv.
`.to_dict()`	Convert to a plain dictionary for JSON reports.
`score_and_select(stats, layers, prune_layers, experts_per_layer, allow_unobserved)`	Rank experts independently in every layer; select in `prune_layers` only.

Connections

Uses: `moe_pruning.model.architecture`, `moe_pruning.scoring.easyep`, `moe_pruning.scoring.statistics`

Used by: `scripts/score_experts.py`