manifest.py

`src/moe_pruning/pruning/manifest.py · 69 lines · module moe_pruning.pruning.manifest`

What it does

Builds the pruning manifest: the single description of what was replaced and how (model, method, layers, expert ids, state-dict prefixes, bias initialisation details, calibration fingerprint).

How it works

- Written into every checkpoint; the loader uses it to rebuild the pruned model and the scripts use it to check that artifacts belong together.

Defined in this file

Extracted from the source code (classes, public functions and methods).

Name	Description
`build_pruning_manifest(cfg, adapter, model_type, layers, selection, bias_init, replacement_report, calibration_fingerprint, stage)`	Assemble the manifest describing the replacement.

Settings it reads

`BIAS_INITIALIZATION`, `HF_MODEL_ID`, `MODEL_DTYPE`, `MODEL_PATH or HF_MODEL_ID`, `PRUNE_EXPERTS_PER_LAYER`, `PRUNE_LAYERS`, `PRUNE_UNOBSERVED_POLICY`, `PRUNING_METHOD`, `REPLACEMENT_METHOD`

Connections

Uses: `moe_pruning.config`, `moe_pruning.model.architecture`

Used by: `scripts/replace_experts.py`