build_c4_support_eval.py — technical

`scripts/build_c4_support_eval.py · 160 lines · scripts/build_c4_support_eval.py`

Purpose

Streams C4 and writes the filtered support evaluation set.

Step by step

Classes and functions in source order. Signatures, line ranges, raised exceptions and settings are extracted from the code; numbered steps describe what the code does.

`main`

```
def main() -> None
```

lines 37-156 · raises SystemExit · settings C4_CONFIG, C4_DATASET_ID, C4_DOMAIN, C4_MAX_SAMPLES, C4_MAX_SCANNED, C4_MAX_TEXT_LENGTH, C4_MIN_KEYWORD_MATCHES, C4_MIN_TEXT_LENGTH, C4_SPLIT, EVALUATION_DATASET_PATH, SEED

Script entry point: runs this stage end to end.

1. Build `SupportDocumentFilter` from the `C4_*` settings; target = `--preview N` or `C4_MAX_SAMPLES`.
2. `load_dataset(C4_DATASET_ID, C4_CONFIG, split=C4_SPLIT, streaming=True)` - rows are fetched shard by shard.
3. For each row until the target or `C4_MAX_SCANNED` is reached: evaluate the filter (rejections counted by reason), drop duplicates, keep `text`, `url`, `timestamp`, score, matched keywords and the C4 row index.
4. Log counts and the top rejection reasons.
5. Preview mode: print shortened accepted documents and stop.
6. Otherwise: error if nothing accepted; warn (with acceptance rate) if fewer than requested; write the JSONL and a manifest (dataset, filter settings, counts, score min/median/max, limitation note).

Errors raised

Function	Exception	Message (f-string fields in braces)
`main`	SystemExit	No documents matched the support filter. Increase C4_MAX_SCANNED, lower C4_MIN_KEYWORD_MATCHES or extend C4_SUPPORT_KEYWORDS.