manifest.py — technical

`src/moe_pruning/pruning/manifest.py · 69 lines · moe_pruning.pruning.manifest`

Purpose

Manifest assembly.

Step by step

Classes and functions in source order. Signatures, line ranges, raised exceptions and settings are extracted from the code; numbered steps describe what the code does.

`build_pruning_manifest`

```
def build_pruning_manifest(cfg: ExperimentConfig, adapter: MoEAdapter, model_type: str, layers: dict[str, MoELayerInfo], selection: dict[str, list[int]], bias_init: dict[str, dict[int, dict]], replacement_report: dict, calibration_fingerprint: str, stage: str) -> dict
```

lines 20-69 · settings BIAS_INITIALIZATION, HF_MODEL_ID, MODEL_DTYPE, MODEL_PATH or HF_MODEL_ID, PRUNE_EXPERTS_PER_LAYER, PRUNE_LAYERS, PRUNE_UNOBSERVED_POLICY, PRUNING_METHOD, REPLACEMENT_METHOD

Assemble the manifest describing the replacement.

1. Top level: stage, timestamp, model source/id/revision/type, dtype, selection settings, replacement method and description, bias initialisation (bias only), calibration fingerprint, replacement report.
2. Per layer: module path, stack, method, number of experts, sorted removed ids, state-dict prefixes per expert, per-expert initialisation results (bias only).

Module constants

Name	Value
`REPLACEMENT_DESCRIPTIONS`	{'bias': 'constant vector b_e; the parent MoE layer multiplies it by the router gate exactly once (g*b)', 'dro...